

```
In [2]: #exp:7
```

```
In [3]: #Name:Vedant M.Padole  
#Roll no:42  
#Sec:C  
#Subject:ET1  
#Date:
```

```
In [4]: import pandas as pd
```

```
In [6]: import os
```

```
In [7]: os.getcwd()
```

```
Out[7]: 'C:\\\\Users\\\\DELL'
```

```
In [8]: os.chdir('C:\\\\Users\\\\DELL\\\\Desktop')
```

```
In [9]: df=pd.read_csv("Salary_Data.csv")
```

```
In [10]: df.head()
```

```
Out[10]: YearsExperience Salary
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
In [11]: df.tail()
```

```
Out[11]: YearsExperience Salary
```

	YearsExperience	Salary
2	9.0	105582.0
5	9.5	116969.0
2	9.6	112635.0
6	10.3	122391.0
2	10.5	121872.0
7		

```
In [12]: df.shape
```

```
Out[12]: (30, 2)
```

```
2
```

```
In [13]: df.size
```

```
9
```

```
Out[13]: 60
```

```
In [14]: df.ndim
```

```
Out[14]: 2
```

```
In [15]: df.describe
```

```
Out[15]: <bound method NDFrame.describe of      YearsExperience      Salary
 0           1.1    39343.0
 1           1.3    46205.0
 2           1.5    37731.0
 3           2.0    43525.0
 4           2.2    39891.0
 5           2.9    56642.0
 6           3.0    60150.0
 7           3.2    54445.0
 8           3.2    64445.0
 9           3.7    57189.0
 10          3.9    63218.0
 11          4.0    55794.0
 12          4.0    56957.0
 13          4.1    57081.0
 14          4.5    61111.0
 15          4.9    67938.0
 16          5.1    66029.0
 17          5.3    83088.0
 18          5.9    81363.0
 19          6.0    93940.0
 20          6.8    91738.0
 21          7.1    98273.0
 22          7.9   101302.0
 23          8.2   113812.0
 24          8.7   109431.0
 25          9.0   105582.0
 26          9.5   116969.0
 27          9.6   112635.0
 28         10.3  122391.0
 29         10.5  121872.0>
```

```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column            Non-Null Count Dtype  
 --- 
 0   YearsExperience  30 non-null      float64 
 1   Salary            30 non-null      float64 
dtypes: float64(2)
memory usage: 608.0 bytes
```

```
In [17]: df.isna()
```

```
Out[17]:  YearsExperience  Salary
```

	YearsExperience	Salary
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	False
7	False	False
8	False	False

	YearsExperience	Salary
9	False	False
10	False	False
11	False	False
12	False	False
13	False	False
14	False	False
15	False	False
16	False	False
17	False	False
18	False	False
19	False	False
20	False	False
21	False	False
22	False	False
23	False	False
24	False	False
25	False	False
26	False	False
27	False	False
28	False	False
29	False	False
	False	False

```
In [18]: df.isna().any()
```

```
Out[18]: YearsExperience    False
          Salary            False
          dtype: bool
```

```
In [19]: df.isna().sum()
```

```
Out[19]: YearsExperience    0
          Salary            0
          dtype: int64
```

```
In [20]: x=df.drop('Salary',axis=1)
```

```
In [21]: x.head()
```

```
Out[21]: YearsExperience
```

0	1.1
1	1.3
2	1.5
3	2.0

YearsExperience

```
4          2.2
```

```
In [22]: y=df.Salary
```

```
In [23]: y.head()
```

```
Out[23]: 0    39343.0
1    46205.0
2    37731.0
3    43525.0
4    39891.0
Name: Salary, dtype: float64
```

```
In [24]: from sklearn.model_selection import train_test_split
```

```
In [25]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_stat
```

```
In [26]: print(x_train.shape)
```

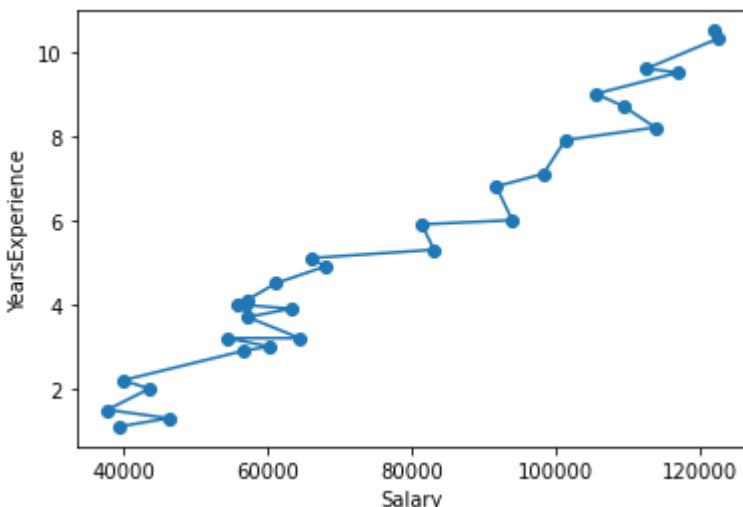
```
(24, 1)
```

```
In [27]: x_test.shape
```

```
Out[27]: (6, 1)
```

```
In [28]: import matplotlib.pyplot as plt
```

```
In [29]: plt.plot(df['Salary'], df['YearsExperience'], marker='o')
plt.xlabel("Salary")
plt.ylabel("YearsExperience")
plt.show()
```



Model Fitting

```
In [32]: from sklearn.linear_model import LinearRegression
LR=LinearRegression()
LR.fit(x_train,y_train)
```

```
Out[32]: LinearRegression()
```

```
In [33]: #Assigning Coefficient (Slope) to m  
m=LR.coef_
```

```
In [34]: print("Coefficient :" , m)  
Coefficient : [9312.57512673]
```

```
In [35]: #Assigning Y-intercept to a  
c=LR.intercept_
```

```
In [36]: print("Intercept : " , c)  
Intercept : 26780.099150628186
```

Evaluation Metrics

```
In [37]: from sklearn import metrics
```

```
In [38]: Accuracy = LR.score(x_test, y_test)  
Accuracy
```

```
Out[38]: 0.988169515729126
```

Conclusion :

The simple linear regression model shows a clear linear relationship between the variables, allowing us to predict outcomes based on this trend with reasonable accuracy

```
In [ ]:
```